

Fundamentals of Neural Networks

Mathias Jackermeier

June 12, 2018

Technische Universität München



Figure 1: A self-driving car.

Credit: Marc van der Chijs / CC BY-ND 2.0

Introduction



Figure 2: A digital assistant.

Credit: Kārlis Dambrāns / CC BY 2.0

Outline

The Perceptron

MNIST Data Sample




Figure 3: Examples from the MNIST database.

Credit: Josef Steppan / CC BY-SA 4.0

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1
- **Idea:** Assign a weight to every input pixel

The perceptron accepts n input values and computes an output value \hat{y} :

$$\begin{aligned}\hat{y} &= \text{sign} \left(\sum_{i=1}^n w_i x_i \right) \\ &\equiv \hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} \right)\end{aligned}\tag{1}$$

Visual Representation




Figure 4: A visual representation of the perceptron model.

Generalizations

- The perceptron is often used in a modified form

Generalizations

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

Generalizations

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (3)$$

Generalizations

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (3)$$

- These modified perceptrons are often called *neurons* or simply *units*

Generalizations

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (3)$$

- These modified perceptrons are often called *neurons* or simply *units*
- **Notation:** We denote the *weighted input* as

$$z = \mathbf{w}^\top \mathbf{x} + b \quad (4)$$

Feedforward Neural Networks

- **Idea:** A combination of multiple neurons could make much better predictions

- **Idea:** A combination of multiple neurons could make much better predictions
- A feedforward neural network is a layered architecture of neurons

Visual Representation




Figure 5: A three-layer feedforward neural network.

- The design of the output layer depends on the task that we wish to perform

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron

The Logistic Sigmoid Function




Figure 6: The logistic sigmoid function $\sigma(x) = \frac{1}{1+\exp(-x)}$.

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron
- *Multiclass classification*: k output units with the softmax function

$$\text{softmax}(x) = \frac{\exp(x)}{\sum_{i=1}^k \exp(z_i)} \quad (5)$$

- The task does not give us any information about how to design the hidden layers

Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function

The Rectified Linear Function




Figure 7: The rectified linear function $g(x) = \max\{0, x\}$.

Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function
- Experimentation and trial & error

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l
- $f^{(l)}$ is the activation function used in the l^{th} layer

- The output at layer l is then given by

$$\mathbf{a}^{(l)} = f^{(l)} \left(\mathbf{W}^{(l)\top} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \right) \quad (6)$$

Training Feedforward Neural Networks

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}
- **Idea:** Iteratively adjust the parameters of the neural network

- The cost function $J(\theta)$ is a measure of how good the network performs

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function
- The total cost is a sum over the costs of the individual training examples:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}, \theta) \quad (7)$$

- In regression, the per-example loss is commonly

$$\mathcal{L}(\mathbf{x}, y, \theta) = \frac{1}{2}(\hat{y} - y)^2 \quad (8)$$

- In binary classification, we often use the cross-entropy loss

$$\mathcal{L}(\mathbf{x}, y, \boldsymbol{\theta}) = -y \ln \hat{y} - (1 - y) \ln(1 - \hat{y}) \quad (9)$$

- In multiclass classification, the cross-entropy becomes

$$\mathcal{L}(\mathbf{x}, \mathbf{y}, \theta) = -\ln \hat{y}_i \quad (10)$$

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- To minimize $J(\theta)$, make small updates in the negative direction of the gradient:

$$\Delta\theta = -\eta\nabla J(\theta) \tag{11}$$

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- To minimize $J(\theta)$, make small updates in the negative direction of the gradient:

$$\Delta\theta = -\eta\nabla J(\theta) \quad (11)$$

- *Stochastic* Gradient Descent computes only an approximation of the gradient

Stochastic Gradient Descent




Figure 8: Stochastic Gradient Descent.

Created with <https://academo.org/demos/3d-surface-plotter/>

Back-propagation

- The back-propagation algorithm efficiently computes the gradient of the cost function

Back-propagation

- The back-propagation algorithm efficiently computes the gradient of the cost function
- It can be derived by recursively applying the chain rule to the layers of the neural network, beginning with the output layer

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient
6. Repeat until the cost is low enough

Conclusion

Thank you!