

Fundamentals of Neural Networks

Mathias Jackermeier

June 12, 2018

Technische Universität München

2018-06-11

Fundamentals of Neural Networks

Fundamentals of Neural Networks

Mathias Jackermeier

June 12, 2018

Technische Universität München



Figure 1: A self-driving car.

Credit: Marc van der Chijs / CC BY-ND 2.0



Figure 1: A self-driving car.
Credit: Marc van der Chijs / CC BY-ND 2.0



Figure 2: A digital assistant.
Credit: Kārlis Dambrāns / CC BY 2.0




Figure 2: A digital assistant.
Credit: Kārlis Dambrāns / CC BY 2.0

Jedes intelligente System benutzt neuronale Netze

2018-06-11

Fundamentals of Neural Networks

└ The Perceptron

The Perceptron

The Perceptron

- Predict whether an input image of a handwritten digit shows a zero or another digit




Figure 3: Examples from the MNIST database.
Credit: Josef Steppan / CC BY-SA 4.0

2018-06-11

Fundamentals of Neural Networks

└ The Perceptron

└ MNIST Data Sample

MNIST Data Sample



Figure 3: Examples from the MNIST database.
Credit: Josef Steppan / CC BY-SA 4.0

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$

Example Task

Fundamentals of Neural Networks

└ The Perceptron

└ Example Task

2018-06-11

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1
- **Idea:** Assign a weight to every input pixel

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1
- **Idea:** Assign a weight to every input pixel

The perceptron accepts n input values and computes an output value \hat{y} :

$$\begin{aligned}\hat{y} &= \text{sign} \left(\sum_{i=1}^n w_i x_i \right) \\ &\equiv \hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} \right)\end{aligned}\quad (1)$$

The perceptron accepts n input values and computes an output value \hat{y} :

$$\begin{aligned}\hat{y} &= \text{sign} \left(\sum_{i=1}^n w_i x_i \right) \\ &\equiv \hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} \right)\end{aligned}\quad (1)$$



Figure 4: A visual representation of the perceptron model.




Figure 4: A visual representation of the perceptron model.

1950

Neuron im Gehirn: Nervenzelle Eingaben Ausgaben

- The perceptron is often used in a modified form

2018-06-11

Fundamentals of Neural Networks

└ The Perceptron

└ Generalizations

- The perceptron is often used in a modified form

f: Aktivierungsfunktion

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$

Generalizations

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$

Fundamentals of Neural Networks

└ The Perceptron

└ Generalizations

f: Aktivierungsfunktion

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$

└ The Perceptron

└ Generalizations

f: Aktivierungsfunktion

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:
$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$
- The sign function can be replaced with a generic function f :
$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$

- These modified perceptrons are often called *neurons* or simply *units*

└ The Perceptron

└ Generalizations

f: Aktivierungsfunktion

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:
$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$
- The sign function can be replaced with a generic function f :
$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$
- These modified perceptrons are often called *neurons* or simply *units*

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:

$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$

- The sign function can be replaced with a generic function f :

$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$

- These modified perceptrons are often called *neurons* or simply *units*
- **Notation:** We denote the *weighted input* as

$$z = \mathbf{w}^\top \mathbf{x} + b \quad (4)$$

└ The Perceptron

└ Generalizations

f: Aktivierungsfunktion

- The perceptron is often used in a modified form
- A scalar bias value can be added to the output computation:
$$\hat{y} = \text{sign}(\mathbf{w}^\top \mathbf{x} + b) \quad (2)$$
- The sign function can be replaced with a generic function f :
$$\hat{y} = f(\mathbf{w}^\top \mathbf{x} + b) \quad (3)$$
- These modified perceptrons are often called *neurons* or simply *units*
- **Notation:** We denote the *weighted input* as
$$z = \mathbf{w}^\top \mathbf{x} + b \quad (4)$$

Feedforward Neural Networks

Perzeptron linear: Viele Funktionen können nicht gelernt werden

Networks of neurons

2018-06-11

Fundamentals of Neural Networks

└─ Feedforward Neural Networks

└─ Networks of neurons

- **Idea:** A combination of multiple neurons could make much better predictions

- **Idea:** A combination of multiple neurons could make much better predictions

Angelehnt an Gehirn

Auch Multilayer Perzeptron

Networks of neurons

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Networks of neurons

- **Idea:** A combination of multiple neurons could make much better predictions
- A feedforward neural network is a layered architecture of neurons

- **Idea:** A combination of multiple neurons could make much better predictions
- A feedforward neural network is a layered architecture of neurons

Angelehnt an Gehirn

Auch Multilayer Perzeptron

Networks of neurons

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Networks of neurons

- **Idea:** A combination of multiple neurons could make much better predictions
- A feedforward neural network is a layered architecture of neurons
- The input of a layer is the output of the previous layer

- **Idea:** A combination of multiple neurons could make much better predictions
- A feedforward neural network is a layered architecture of neurons
- The input of a layer is the output of the previous layer

Angelehnt an Gehirn

Auch Multilayer Perzeptron

Visual Representation




Figure 5: A three-layer feedforward neural network.

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Visual Representation




Figure 5: A three-layer feedforward neural network.

Tiefe

Verbindungen nur zwischen Layers

- The design of the output layer depends on the task that we wish to perform

2018-06-11

Fundamentals of Neural Networks

└─ Feedforward Neural Networks

└─ Output Layer

Output Layer

- The design of the output layer depends on the task that we wish to perform

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Output Layer

Output Layer

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron

Output Layer

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Output Layer

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron

The Logistic Sigmoid Function




Figure 6: The logistic sigmoid function $\sigma(x) = \frac{1}{1+\exp(-x)}$

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ The Logistic Sigmoid Function

The Logistic Sigmoid Function




Figure 6: The logistic sigmoid function $\sigma(x) = \frac{1}{1+\exp(-x)}$

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron
- *Multiclass classification*: k output units with the softmax function

$$\text{softmax}(x) = \frac{\exp(x)}{\sum_{i=1}^k \exp(z_i)} \quad (5)$$

- The design of the output layer depends on the task that we wish to perform
- *Regression*: one single linear neuron
- *Binary classification*: one single sigmoid neuron
- *Multiclass classification*: k output units with the softmax function

$$\text{softmax}(x) = \frac{\exp(x)}{\sum_{i=1}^k \exp(z_i)} \quad (5)$$

softmax: Generalisierung, normalisiert

Viele andere Probleme: Auffälliges Verhalten, Bildgenerierung

Hidden Layers

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Hidden Layers

- The task does not give us any information about how to design the hidden layers

- The task does not give us any information about how to design the hidden layers

Daher kommt Deep Learning
Verschiedene Repräsentationen (Kanten, Objekte,...) \Rightarrow Abstraktionen!
Kontinuierlich, nicht linear!

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view

Hidden Layers

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view

Daher kommt Deep Learning

Verschiedene Repräsentationen (Kanten, Objekte,...) \Rightarrow Abstraktionen!

Kontinuierlich, nicht linear!

Hidden Layers

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice

Daher kommt Deep Learning

Verschiedene Repräsentationen (Kanten, Objekte,...) \Rightarrow Abstraktionen!

Kontinuierlich, nicht linear!

Hidden Layers

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function

Daher kommt Deep Learning

Verschiedene Repräsentationen (Kanten, Objekte,...) \Rightarrow Abstraktionen!

Kontinuierlich, nicht linear!

The Rectified Linear Function




Figure 7: The rectified linear function $g(x) = \max\{0, x\}$

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ The Rectified Linear Function

The Rectified Linear Function




Figure 7: The rectified linear function $g(x) = \max\{0, x\}$

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function
- Experimentation and trial & error

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Hidden Layers

- The task does not give us any information about how to design the hidden layers
- Depth: Irrelevant from a theoretical point of view
- Deep networks perform almost always better in practice
- Activation function: Three common choices are the logistic sigmoid, the tanh, and the rectified linear function
- Experimentation and trial & error

Mathematical Formulation

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b

f wird komponentenweise angewendet

Mathematical Formulation

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight matrices $\mathbf{W}^{(l)}$ and bias vectors $\mathbf{b}^{(l)}$ to specify the parameters of a layer l

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l

f wird komponentenweise angewendet

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight matrices $\mathbf{W}^{(l)}$ and bias vectors $\mathbf{b}^{(l)}$ to specify the parameters of a layer l
- $f^{(l)}$ is the activation function used in the l^{th} layer

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l
- $f^{(l)}$ is the activation function used in the l^{th} layer

f wird komponentenweise angewendet

Mathematical Formulation

2018-06-11

Fundamentals of Neural Networks

└ Feedforward Neural Networks

└ Mathematical Formulation

- The output at layer l is then given by

$$\mathbf{a}^{(l)} = f^{(l)} \left(\mathbf{W}^{(l)\top} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \right) \quad (6)$$

$$\mathbf{a}^0 = \mathbf{x}$$

$$\hat{y} = \mathbf{a}^L$$

- The output at layer l is then given by

$$\mathbf{a}^{(l)} = f^{(l)} \left(\mathbf{W}^{(l)\top} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \right) \quad (6)$$

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

Training Feedforward Neural
Networks

Training Feedforward Neural Networks

Training Scenario

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Training Scenario

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)})$ with corresponding labels \mathbb{Y}

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}

Zufällig initialisiert

Training Scenario

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Training Scenario

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}

Zufällig initialisiert

Training Scenario

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Training Scenario

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}
- **Idea:** Iteratively adjust the parameters of the neural network

- We have training examples $\mathbb{X} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)})$ with corresponding labels \mathbb{Y}
- We want to learn a mapping from \mathbb{X} to \mathbb{Y}
- **Idea:** Iteratively adjust the parameters of the neural network

Zufällig initialisiert

Cost Functions

- The cost function $J(\theta)$ is a measure of how good the network performs

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Cost Functions

Von den Parametern zu einem Skalar

Größer als 0

Auch Loss oder Error

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function

Cost Functions

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Cost Functions

Von den Parametern zu einem Skalar

Größer als 0

Auch Loss oder Error

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function
- The total cost is a sum over the costs of the individual training examples:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}, \theta) \quad (7)$$

Von den Parametern zu einem Skalar

Größer als 0

Auch Loss oder Error

- The cost function $J(\theta)$ is a measure of how good the network performs
- Learning can be framed as minimizing the cost function
- The total cost is a sum over the costs of the individual training examples:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\mathbf{x}^{(i)}, \mathbf{y}^{(i)}, \theta) \quad (7)$$

Mean squared error

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Mean squared error

- In regression, the per-example loss is commonly

$$\mathcal{L}(\mathbf{x}, y, \theta) = \frac{1}{2}(\hat{y} - y)^2 \quad (8)$$

- In regression, the per-example loss is commonly

$$\mathcal{L}(\mathbf{x}, y, \theta) = \frac{1}{2}(\hat{y} - y)^2 \quad (8)$$

Label: Skalar was wir vorhersagen wollen

Distanz

Erfüllt Bedingungen

- In binary classification, we often use the cross-entropy loss

$$\mathcal{L}(\mathbf{x}, y, \theta) = -y \ln \hat{y} - (1 - y) \ln(1 - \hat{y}) \quad (9)$$

MSE schlecht in Klassifikation

Label: 1 oder 0

- In multiclass classification, the cross-entropy becomes

$$\mathcal{L}(\mathbf{x}, \mathbf{y}, \theta) = -\ln \hat{y}_i \quad (10)$$

Label: i-te Klasse

Maximum Likelihood Estimation

Stochastic Gradient Descent

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Stochastic Gradient Descent

⇒ Kleine Änderungen in die entgegengesetzte Richtung des Gradienten
Learning Rate Hyperparameter durch rumexperimentieren
Erweiterungen

Stochastic Gradient Descent

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Stochastic Gradient Descent

⇒ Kleine Änderungen in die entgegengesetzte Richtung des Gradienten
Learning Rate Hyperparameter durch rumexperimentieren
Erweiterungen

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$

Stochastic Gradient Descent

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$

- To minimize $J(\theta)$, choose

$$\Delta\theta = -\eta \nabla J(\theta), \quad (12)$$

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Stochastic Gradient Descent

2018-06-11

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$

- To minimize $J(\theta)$, choose

$$\Delta\theta = -\eta \nabla J(\theta), \quad (12)$$

⇒ Kleine Änderungen in die entgegengesetzte Richtung des Gradienten
Learning Rate Hyperparameter durch rumexperimentieren
Erweiterungen

Stochastic Gradient Descent

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$

- To minimize $J(\theta)$, choose

$$\Delta\theta = -\eta \nabla J(\theta), \quad (12)$$

- *Stochastic* Gradient Descent computes only an approximation of the gradient

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Stochastic Gradient Descent

2018-06-11

- (Stochastic) Gradient Descent is the most common algorithm to minimize cost functions in neural networks
- A change $\Delta\theta$ in the parameters corresponds roughly to the change

$$\Delta J(\theta) \approx \nabla J(\theta)^\top \Delta\theta \quad (11)$$
- To minimize $J(\theta)$, choose

$$\Delta\theta = -\eta \nabla J(\theta), \quad (12)$$
- Stochastic Gradient Descent computes only an approximation of the gradient

⇒ Kleine Änderungen in die entgegengesetzte Richtung des Gradienten
Learning Rate Hyperparameter durch rumexperimentieren
Erweiterungen




Figure 8: Stochastic Gradient Descent.

Created with <https://academo.org/demos/3d-surface-plotter/>

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Stochastic Gradient Descent




Figure 8: Stochastic Gradient Descent.
Created with <https://academo.org/demos/3d-surface-plotter/>

Back-propagation

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

- The back-propagation algorithm efficiently computes the gradient of the cost function

- The back-propagation algorithm efficiently computes the gradient of the cost function

Back-propagation

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

- The back-propagation algorithm efficiently computes the gradient of the cost function
- It can be derived by recursively applying the chain rule to the layers of the neural network, beginning with the output layer

- The back-propagation algorithm efficiently computes the gradient of the cost function
- It can be derived by recursively applying the chain rule to the layers of the neural network, beginning with the output layer

Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation





Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation



Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back propagation




Figure 9: The Back-propagation algorithm.

Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back-propagation




Figure 9: The Back-propagation algorithm.

Back-propagation





Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation



Back-propagation




Figure 9: The Back-propagation algorithm.

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ Back-propagation

Back-propagation




Figure 9: The Back-propagation algorithm.

1. Propagate all training examples of a minibatch forward through the network

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient

The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient
6. Repeat until the cost is low enough

2018-06-11

Fundamentals of Neural Networks

└ Training Feedforward Neural Networks

└ The Complete Learning Algorithm

1. Propagate all training examples of a minibatch forward through the network
2. Compute the cost for each training example
3. Compute all gradients using back-propagation
4. Compute the average gradient
5. Update the parameters in the negative direction of the gradient
6. Repeat until the cost is low enough

Conclusion

2018-06-11

Fundamentals of Neural Networks

└ Conclusion

Conclusion

Komplexe Netzwerke einfacher Einheiten

Abstraktionen

Lernen: Kleine Updates der Parameter so dass das Netzwerk besser wird

Überall in Deep Learning

Viele weitere Anwendungen in der Zukunft

Thank you!