

Fundamentals of Neural Networks

Mathias Jackermeier

June 6, 2018

Technische Universität München



Figure 1: A self-driving car.

Credit: Marc van der Chijs / CC BY-ND 2.0

Introduction



Figure 2: A digital assistant.

Credit: Kārlis Dambrāns / CC BY 2.0

Introduction



Airplane



Car



Person

Figure 3: Object detection in images.

Credit: Lu et. al¹

¹"1-HKUST: Object Detection in ILSVRC 2014", *CoRR*, vol. abs/1409.6155, 2014

Outline

1. The Perceptron
2. Feedforward Neural Networks
 - Architecture
 - Mathematical formulation
3. Training Feedforward Neural Networks
 - Cost functions
 - Stochastic Gradient Descent
 - Back-propagation
4. Extensions

The Perceptron

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit

MNIST Data Sample

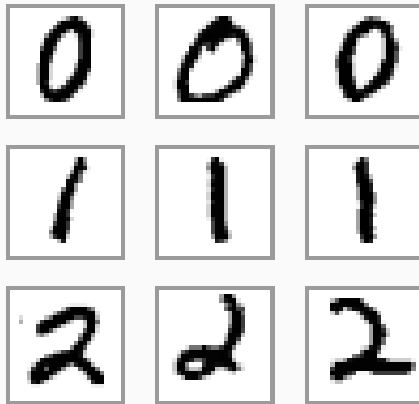


Figure 4: Examples from the MNIST database.
Credit: Josef Steppan / CC BY-SA 4.0

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1

Example Task

- Predict whether an input image of a handwritten digit shows a zero or another digit
- The image is represented as a flattened vector of pixel intensities $\mathbf{x} \in \mathbb{R}^{784}$
- The output should be 1 if the image shows a zero, otherwise it should be -1
- **Idea:** Assign a weight to every input pixel

The perceptron accepts n input values and computes an output value \hat{y} :

$$\begin{aligned}\hat{y} &= \text{sign} \left(\sum_{i=1}^n w_i x_i \right) \\ &\equiv \hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} \right)\end{aligned}\tag{1}$$

Visual Representation

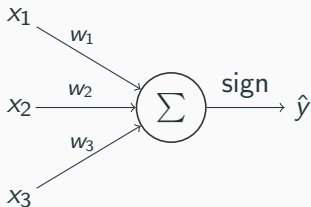


Figure 5: A visual representation of the perceptron model.

Shortcomings of the Perceptron

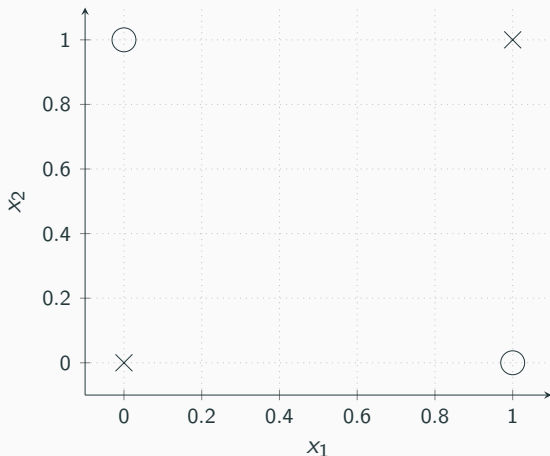


Figure 6: The perceptron cannot learn the XOR function since the data is not linearly separable.

Feedforward Neural Networks

- **Idea:** A combination of multiple perceptrons could make much better predictions

Networks of perceptron-like units

- **Idea:** A combination of multiple perceptrons could make much better predictions
- We arrange the perceptrons in layers

Networks of perceptron-like units

- **Idea:** A combination of multiple perceptrons could make much better predictions
- We arrange the perceptrons in layers
- The input of a layer is the output of the previous layer

Networks of perceptron-like units

- **Idea:** A combination of multiple perceptrons could make much better predictions
- We arrange the perceptrons in layers
- The input of a layer is the output of the previous layer
- This network model is called *feedforward neural network* or *multilayer perceptron*

Visual Representation




Figure 7: A three-layer feedforward neural network.

- In practice, we use a modified version of the perceptron as computing units

- In practice, we use a modified version of the perceptron as computing units
- We add a scalar bias value to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- In practice, we use a modified version of the perceptron as computing units
- We add a scalar bias value to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- We replace the sign function with a generic function f :

$$\hat{y} = f \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (3)$$

- In practice, we use a modified version of the perceptron as computing units
- We add a scalar bias value to the output computation:

$$\hat{y} = \text{sign} \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (2)$$

- We replace the sign function with a generic function f :

$$\hat{y} = f \left(\mathbf{w}^\top \mathbf{x} + b \right) \quad (3)$$

- These modified perceptrons are often called *neurons* or simply *units*

Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b

Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l

Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l
- The weight $w_{ij}^{(l)}$ is the weight from the i th neuron in the $l-1$ th layer to the j th neuron in the l th layer

Mathematical Formulation

- We can specify a single neuron with a weight vector \mathbf{w} and a bias value b
- Since a neural network consists of multiple neurons in a layer, we need weight *matrices* $\mathbf{W}^{(l)}$ and bias *vectors* $\mathbf{b}^{(l)}$ to specify the parameters of a layer l
- The weight $w_{ij}^{(l)}$ is the weight from the j neuron in the $l-1$ layer to the i neuron in the l layer
- The bias $b_i^{(l)}$ is the bias of the i neuron in the l layer

- The output at layer l is then given by

$$\mathbf{a}^{(l)} = f^{(l)} \left(\mathbf{W}^{(l)\top} \mathbf{a}^{(l-1)} + \mathbf{b}^{(l)} \right) \quad (4)$$

Training Feedforward Neural Networks

Extensions

Thank you!